

- Resource-Efficient Real-Time Video Analytics. In *Proceedings of the Annual conference of the ACM Special Interest Group on Data Communication on the applications, technologies, architectures, and protocols for computer communication*, 359–376.
- Liu, S.; Lin, Y.; Zhou, Z.; Nan, K.; Liu, H.; and Du, J. 2018. On-demand deep model compression for mobile devices: A usage-driven model selection framework. In *Proceedings of the 16th Annual International Conference on Mobile Systems, Applications, and Services*, 389–400.
- Marszalek, M.; Laptev, I.; and Schmid, C. 2009. Actions in context. In *2009 IEEE Conference on Computer Vision and Pattern Recognition*, 2929–2936. IEEE.
- MindSpore. 2021. MindSopre. <https://github.com/mindspore-ai/mindspore>. Accessed: 2021-08-01.
- Mozilla. 2021. DeepSpeech. <https://github.com/mozilla/DeepSpeech>. Accessed: 2021-08-01.
- Ning, L.; Guan, H.; and Shen, X. 2019. Adaptive Deep Reuse: Accelerating CNN Training on the Fly. In *2019 IEEE 35th International Conference on Data Engineering (ICDE)*, 1538–1549.
- Olafenwa, M. 2021a. Action-Net. <https://github.com/OlafenwaMoses/Action-Net>. Accessed: 2021-08-01.
- Olafenwa, M. 2021b. Traffic-Net. <https://github.com/OlafenwaMoses/Traffic-Net>. Accessed: 2021-08-01.
- Passalis, N.; Tzelepi, M.; and Tefas, A. 2020. Heterogeneous knowledge distillation using information flow modeling. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2339–2348.
- PyTorch. 2021. PyTorch. <https://github.com/pytorch/pytorch>. Accessed: 2021-08-01.
- Redmon, J.; and Farhadi, A. 2018. Yolov3: An incremental improvement. *arXiv preprint arXiv:1804.02767*.
- Sanh, V.; Wolf, T.; and Ruder, S. 2019. A hierarchical multi-task approach for learning embeddings from semantic tasks. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 6949–6956.
- Selvaraju, R. R.; Cogswell, M.; Das, A.; Vedantam, R.; Parikh, D.; and Batra, D. 2020. Grad-CAM: Visual Explanations from Deep Networks via Gradient-based Localization. *International Journal of Computer Vision*, 128(2): 336–359.
- Serengil, S. I.; and Ozpinar, A. 2020. LightFace: A Hybrid Deep Face Recognition Framework. In *2020 Innovations in Intelligent Systems and Applications Conference (ASYU)*, 23–27. IEEE.
- Shen, Z.; He, Z.; and Xue, X. 2019. Meal: Multi-model ensemble via adversarial learning. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 4886–4893.
- Song, C.; and Shmatikov, V. 2020. Overlearning Reveals Sensitive Attributes. In *International Conference on Learning Representations*.
- Sutskever, I.; Vinyals, O.; and Le, Q. V. 2014. Sequence to sequence learning with neural networks. In *Advances in neural information processing systems*, 3104–3112.
- Sviridenko, M. 2004. A note on maximizing a submodular set function subject to a knapsack constraint. *Operations Research Letters*, 32(1): 41–43.
- Tan, C.; Sun, F.; Kong, T.; Zhang, W.; Yang, C.; and Liu, C. 2018. A survey on deep transfer learning. In *International conference on artificial neural networks*, 270–279. Springer.
- TensorFlow. 2021. TensorFlow. <https://github.com/tensorflow/tensorflow>. Accessed: 2021-08-01.
- Tieleman, T.; and Hinton, G. 2012. Lecture 6.5-rmsprop: Divide the gradient by a running average of its recent magnitude. *COURSERA: Neural networks for machine learning*, 4(2): 26–31.
- Tran, D.; Bourdev, L.; Fergus, R.; Torresani, L.; and Paluri, M. 2015. Learning spatiotemporal features with 3d convolutional networks. In *Proceedings of the IEEE international conference on computer vision*, 4489–4497.
- Tripuraneni, N.; Jordan, M.; and Jin, C. 2020. On the Theory of Transfer Learning: The Importance of Task Diversity. *Advances in Neural Information Processing Systems*, 33.
- Wang, G. 2021. Image Captioning. https://github.com/DeepRNN/image_captioning. Accessed: 2021-08-01.
- Wang, J.; Fang, Z.; and Zhao, H. 2020. Alignnet: A unifying approach to audio-visual alignment. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, 3309–3317.
- Xie, X.; and Kim, K.-H. 2019. Source compression with bounded dnn perception loss for iot edge computer vision. In *The 25th Annual International Conference on Mobile Computing and Networking*, 1–16.
- Yao, A.; and Sun, D. 2020. Knowledge Transfer via Dense Cross-Layer Mutual-Distillation. In *European Conference on Computer Vision*, 294–311. Springer.
- Ye, H.-J.; Lu, S.; and Zhan, D.-C. 2020. Distilling Cross-Task Knowledge via Relationship Matching. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 12396–12405.
- Yuan, M.; Zhang, L.; Li, X.-Y.; and Xiong, H. 2020a. Comprehensive and efficient data labeling via adaptive model scheduling. In *2020 IEEE 36th International Conference on Data Engineering (ICDE)*, 1858–1861. IEEE.
- Yuan, M.; Zhang, L.; Wu, Z.; and Zheng, D. 2020b. High-quality Activity-Level Video Advertising. In *2020 IEEE/ACM 28th International Symposium on Quality of Service (IWQoS)*, 1–10. IEEE.
- Yuksel, S. E.; Wilson, J. N.; and Gader, P. D. 2012. Twenty years of mixture of experts. *IEEE transactions on neural networks and learning systems*, 23(8): 1177–1193.
- Zhang, Y.; Xiang, T.; Hospedales, T. M.; and Lu, H. 2018. Deep mutual learning. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 4320–4328.
- Zhang, Y.; and Yang, Q. 2021. A survey on multi-task learning. *IEEE Transactions on Knowledge and Data Engineering*.
- Zhou, Z.-H. 2012. *Ensemble methods: foundations and algorithms*. CRC press.